
Workgroup:	Network Working Group
Internet-Draft:	draft-ietf-rpp-core-00
Published:	5 July 2026
Intended Status:	Standards Track
Expires:	6 January 2027
Authors:	M. Wullink P. Kowalik <i>SIDN Labs DENIC</i>

RESTful Provisioning Protocol (RPP)

Abstract

This document describes the endpoints for the RESTful Provisioning Protocol, used for the provisioning and management of objects in a shared database.

Status of This Memo

This Internet-Draft is submitted in full conformance with the provisions of BCP 78 and BCP 79.

Internet-Drafts are working documents of the Internet Engineering Task Force (IETF). Note that other groups may also distribute working documents as Internet-Drafts. The list of current Internet-Drafts is at <https://datatracker.ietf.org/drafts/current/>.

Internet-Drafts are draft documents valid for a maximum of six months and may be updated, replaced, or obsoleted by other documents at any time. It is inappropriate to use Internet-Drafts as reference material or to cite them other than as "work in progress."

This Internet-Draft will expire on 6 January 2027.

Copyright Notice

Copyright (c) 2026 IETF Trust and the persons identified as the document authors. All rights reserved.

This document is subject to BCP 78 and the IETF Trust's Legal Provisions Relating to IETF Documents (<https://trustee.ietf.org/license-info>) in effect on the date of publication of this document. Please review these documents carefully, as they describe your rights and restrictions with respect to this document. Code Components extracted from this document must include Revised BSD License text as described in Section 4.e of the Trust Legal Provisions and are provided without warranty as described in the Revised BSD License.

Table of Contents

1. Introduction	4
2. Terminology	5
3. Conventions Used in This Document	5
4. Mapping to EPP	5
5. Request Headers	6
6. Response Headers	7
7. Error handling and relation between HTTP status codes and RPP codes	8
8. Problem Detail responses for errors	9
9. Bootstrapping	10
10. Discoverability	11
10.1. Workflow	14
11. Versioning	14
11.1. Endpoints	14
11.2. Messages	15
11.3. Extensions	15
11.4. Profiles	15
12. Media types	16
13. Profiles	16
13.1. Definition	16
13.2. Inheritance	17
13.3. Signalling	18
13.3.1. Header signalling	18
13.3.2. Media type parameter signalling	19
13.3.3. Process signalling	19
14. Endpoints	20
14.1. HTTP Mapping Rules	20
14.1.1. Rule 1: Collection Path Segment	20
14.1.2. Rule 2: Uniform Interface Operations	20

14.1.3. Rule 3: Direct Access Sub-Resource Path Segment	21
14.1.4. Rule 4: Process Uniform Interface Operations	22
14.1.5. Rule 5: Extended Process Operations	23
14.1.6. Rule 6: Process Listing	24
14.2. Derived Endpoint Reference	24
14.3. Availability for Creation	26
14.4. Resource Information	27
14.5. Create Resource	28
14.6. Delete Resource	29
14.7. Update Resource	30
14.8. Processes	31
14.8.1. Relation to object representation	32
14.8.2. Response with information about created process	32
14.8.3. Restore Resource	32
14.8.4. Renew Resource	32
14.8.5. Transfer Resource	33
14.8.5.1. Start	33
14.8.5.2. Status	34
14.8.5.3. Cancel	35
14.8.5.4. Reject	36
14.8.5.5. Approve	37
14.9. Messages	38
14.9.1. Retrieve	38
14.9.2. Delete	39
15. Extension Framework	40
16. RPP Result Codes	40
17. Authentication and Authorization	40
18. IANA Considerations	41
18.1. URN Sub-namespace for RPP (urn:ietf:params:rpp)	41
18.2. RPP registry group	41

18.3. RPP Discovery registry	41
18.4. RPP Extension registry	42
18.5. RPP Profile registry	42
18.6. RPP Result Codes Registry	42
18.7. Link Relation Type: rpp-process	43
18.8. RPP Media Type (application/rpp+json)	43
19. Internationalization Considerations	44
20. Security Considerations	44
21. Change History	44
21.1. Version draft-wullink-rpp-core-05 to draft-ietf-rpp-core-00	44
21.2. Version 04 to 05	45
21.3. Version 03 to 04	45
21.4. Version 02 to 03	45
21.5. Version 01 to 02	45
21.6. Version 00 to 01	45
21.7. Version 00 (draft-rpp-core) to 00 (draft-wullink-rpp-core)	45
22. Acknowledgements	46
23. References	46
23.1. Normative References	46
23.2. Informative References	48
Authors' Addresses	48

1. Introduction

This document describes an Application Programming Interface (API) based on the HTTP protocol [\[RFC2616\]](#) and the principles of [\[REST\]](#). Conforming to the REST constraints is generally referred to as being "RESTful". Hence the API is dubbed: "RESTful Provisioning Protocol" or "RPP" for short.

The RPP API is designed to be used for the provisioning and management of objects in a shared database, such as domain names, hosts, and entities.

2. Terminology

In this document the following terminology is used.

REST - Representational State Transfer ([[REST](#)]). An architectural style.

RESTful - A RESTful web service is a web service or API implemented using HTTP and the principles of [[REST](#)].

RESTful Provisioning Protocol or RPP - The protocol described in this document.

URL - A Uniform Resource Locator as defined in [[RFC3986](#)].

Resource - An object having a type, data, and possible relationship to other resources, identified by a URL.

RPP client - An HTTP user agent performing an RPP request

RPP server - An HTTP server responsible for processing requests and returning results in any supported media type.

JWT - JSON Web Token as defined in [[RFC7519](#)].

3. Conventions Used in This Document

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [[RFC2119](#)].

In examples, indentation and white space in examples are provided only to illustrate element relationships and are not REQUIRED features of the protocol.

All example requests assume a RPP server using HTTP version 2 is listening on the standard HTTPS port on host rpp.example. An authorization token has been provided by an out of band process and MUST be used by the client to authenticate each request.

4. Mapping to EPP

RPP is designed as an independent protocol and does not require an EPP server. RPP concepts such as transaction identifiers, result codes, and object attributes are defined in their own right and serve RPP purposes regardless of whether an EPP backend is present, however compatibility with EPP is to the great extent preserved. Implementers with no prior EPP experience are be able to implement RPP based solely on this specification.

Some RPP concepts are functionally similar to EPP concepts, but they are not directly derived from EPP and MAY have different semantics. To avoid confusion, RPP elements SHOULD NOT use an "EPP" prefix or suffix. For implementers who operate an EPP backend and need to bridge RPP requests to EPP commands, a separate RPP-to-EPP mapping document [TODO REF] is provided. Any extensions to RPP are not covered by that mapping document; the mapping of extension elements MUST be defined in the respective extension specification.

5. Request Headers

A RPP request does not always require a request message body. The information conveyed by the HTTP method, URL, and request headers may be sufficient for the server to be able to successfully process a request. However, the client **MUST** include a request message body when the server requires additional attributes to be present in the request message. The RPP HTTP headers listed below use the "RPP-" prefix, following the recommendations in [\[RFC6648\]](#).

- **RPP-Cltrid**: A client-assigned transaction identifier. The client **MUST** include this header in every request. It serves two independent purposes: as an idempotency key, allowing the server to detect and safely handle duplicate requests, and as an audit-trail identifier, enabling end-to-end correlation of a request across client and server logs. The value **MUST** be unique per request.
- **RPP-Authorization**: The client **MAY** use this header to send authorization information in the format `<method> <authorization information>`, similar to the HTTP Authorization header, defined in [\[RFC9110, Section 11.6.2\]](#). The `<method>` indicates the type of authorization being used. For EPP object authorization information, for example the authorization information used for domain names described in [\[RFC5731, Section 2.3\]](#), a new `authinfo` method is defined and **MUST** be used. The `<authorization information>` defines the following comma separated fields:
 - **value** (REQUIRED): Base64 encoded EPP password-based authorization information. Base64 encoding is used to prevent problems when special characters are present that may conflict with the format rules for the Authorization header.
 - **roid** (OPTIONAL): A Roid as defined in [\[RFC5731\]](#), [\[RFC5733\]](#), and [\[RFC5730\]](#). The roid is used to identify the object for which the authorization information is provided. If the roid is not provided, then the server **MUST** assume that the authorization information is linked to the object identified by the URL of the request.

Use of the RPP-Authorization header:

```
RPP-Authorization: authinfo value=TXkgU2VjcmFRva2Vu, roid=REG-X-123
```

The value of the RPP-Authorization header is case sensitive. The server **MUST** reject requests where the case of the header value does not match the expected case. The RPP-Authorization header is specific to the user agent and **MUST NOT** be cached, as recommended by [Section 16.4.2](#), the server **MUST** use the correct HTTP cache directives to prevent caching of the RPP-Authorization header.

- **RPP-Profile**: The client **MUST** use this header to indicate the profiles is used in the request.

6. Response Headers

The server HTTP response contains a status code, headers, and MAY contain an RPP response message in the message body. HTTP headers are used to transmit additional data to the client and MAY be used to send RPP process related data to the client. HTTP headers used by RPP MUST use the "RPP-" prefix, the following response headers have been defined for RPP.

- **RPP-Svtrid**: A server-assigned transaction identifier. The server MUST include this header in every response. It provides a unique, server-side audit-trail reference for the processed request.
- **RPP-Cltrid**: The server MUST echo the client transaction identifier from the request back to the client in this response header. This allows the client to correlate responses to their originating requests.
- **RPP-Code**: This header is the equivalent of the EPP result code defined in [\[RFC5730\]](#) and MUST be used accordingly. This header MUST be added to all responses and MAY be used by the client for easy access to the result code, without having to parse the HTTP response message body.

For the EPP codes related to session management (1500, 2500, 2501 and 2502) there are no corresponding RPP codes.

In order for RPP to be backwards compatible with EPP, RPP will use 5-digit coding of the result codes, where first digit will denote origin specification of the result codes.

For [\[RFC5730\]](#) Result Codes the leading digit MUST be "0". For RPP result codes the leading digit MUST be "1". For avoidance of confusion RPP MUST not define new codes with the same semantic meaning as already defined in EPP.

For RPP codes the remaining 4 digits MUST keep the same semantics as [\[RFC5730\]](#) Result Codes.

- **RPP-Queue-Size**: Return the number of unacknowledged messages in the client message queue. The server MAY include this header in all RPP responses.

When a uniform interface operation implicitly creates a process object as a side effect, the server MUST communicate the URL of the created process resource using the Link response header [\[RFC8288\]](#) with the `rpp-process` relation type. If multiple process objects are created, the server MUST include one Link header field per created process resource, each with `rel="rpp-process"`.

Example:

```
Link: <https://rpp.example/rpp/v1/domainNames/foo.example/processes/
createProcesses/latest>; rel="rpp-process"
```

7. Error handling and relation between HTTP status codes and RPP codes

RPP leverages standard HTTP status codes to reflect the outcome of RPP operations. The RPP result codes are based on the EPP result codes defined in [RFC5730]. This allows clients to handle responses generically using common HTTP patterns. While the HTTP status code provides the primary, high-level outcome, the specific RPP result code MUST still be provided in the RPP-Code HTTP header for detailed diagnostics.

The mapping strategy is to use the most specific HTTP code that accurately reflects the operation's result.

For common and well-defined outcomes, a specific HTTP status code is used. For example, an attempt to access a non-existent resource (EPP code 2302) MUST return 404 Not Found, and an attempt to create a resource that already exists (EPP code 2303) MUST return 409 Conflict. This allows a client to handle these common situations based on the HTTP code alone.

For all other failures, a generic HTTP status code is used. Client-side errors (e.g., syntax, parameter, or policy violations) MUST return 400 Bad Request. Server-side failures MUST return 500 Internal Server Error.

The server MUST return HTTP status codes, following the mapping rules in Table 1.

Table 1: RPP result code and HTTP Status-Code mapping.

HTTP Status-Code	Description	Corresponding RPP result code(s)
Success (2xx)		
200 OK	The request was successful (e.g., for GET or UPDATE).	01000 (in all cases not specified otherwise), 01300, 01301
201 Created	The resource was created successfully.	01000 for resource creating requests (POST/PUT)
202 Accepted	The request was accepted for asynchronous processing.	01001
204 No Content	The resource was deleted successfully.	01000 for DELETE
Client Errors (4xx)		
400 Bad Request	Generic client-side error (syntax, parameters, policy).	02000-02005, 02104-02106, 02300-02301, 02304-02308
403 Forbidden	Authentication or authorization failed.	02200-02202
404 Not Found	The requested resource does not exist.	02303

HTTP Status-Code	Description	Corresponding RPP result code(s)
409 Conflict	The resource could not be created because it already exists.	02302
Server Errors (5xx)		
500 Internal Server Error	Generic server-side error; command failed.	02400
501 Not Implemented	The requested command or feature is not implemented.	02100-02103

Table 1

Some EPP result codes, like 01500, 02500, 02501 and 02502 are related to session management and therefore not applicable to a sessionless RPP protocol.

8. Problem Detail responses for errors

When an error occurs that prevents processing of the requested action, an RPP server MUST respond using a Problem Detail document [RFC9457] detailing what went wrong, or what was not acceptable to the server. The type field MUST be the `urn:ietf:params:rpp:problem` URN. The status field MUST reflect the HTTP status code. The document MUST contain an `errors` element, as a list of objects detailing individual errors.

This document consists of the following fields:

`type` (required, string) This field MUST always contain the fixed string value `urn:ietf:params:rpp:error` to indicate that this is an RPP Problem Detail response.
`title` (required, string) A short human-readable description of the Problem Detail type.
`errors` (optional, array of error objects) MUST contain objects detailing information about the specific errors that occurred, including optional references to which values in the original request were not acceptable to the server.

The error object consist of the following fields:

`type` (required, string) This field SHOULD be a URI that identifies the error type. This URI MAY be dereferenceable to obtain human-readable documentation about the error type.
`result` (required, string) This field MUST contain the RPP result code associated with the error.
`paths` (optional, list of strings) The JSONPath [RFC9535] to the value(s) in the request that caused the error. This field MAY be omitted if no specific value in the request can be identified as the cause of the error.
`reason` (required, string) A human-readable detailed description of the error.

RPP specifications and/or extensions MAY add extension fields to the error object, to convey additional information about the causes of the error. For example, to indicate the account balance on a billing failure, the following Problem Detail response could be used:

```
{
  "type": "urn:ietf:params:rpp:error",
  "title": "RPP Error",
  "errors": [{
    "type": "https://rpp.example/problems/billing/billing-failure",
    "result": "02104",
    "reason": "Not enough balance on account to create domain",
    "balance": 10.0,
    "action_cost": 25.0
  }]
}
```

Problem Detail response containing multiple errors for a domain create request using an invalid domain name (`_.example`), JSONPaths are specified in the error objects to indicate which value in the request caused the error:

```
{
  "type": "urn:ietf:params:rpp:error",
  "title": "RPP Error",
  "errors": [{
    "type": "https://rpp.example/problems/domains/domain-syntax",
    "result": "02005",
    "paths": ["$.domain.name"],
    "reason": "Invalid character(s) in domain name",
  },
  {
    "type": "https://rpp.example/problems/domains/domain-allowed-length",
    "result": "02004",
    "paths": ["$.domain.name"],
    "reason": "Domain name length must be between 3 and 63 characters",
  }
]
```

9. Bootstrapping

The server MUST provide a mechanism for clients to discover the location of the RPP server, two methods are defined for this purpose, using the IANA registry for RPP servers or using DNS-based bootstrapping. The server MUST provide at least one of these methods, and MAY provide both methods. The client MUST use either the IANA registry for RPP servers or a DNS lookup using an HTTPS resource record as defined in [\[RFC9460\]](#). The client MUST be able to handle the case where the chosen method does not return a valid RPP server location, and MUST be able to use the other method as a fallback to discover the RPP server location.

The format and procedure for adding an RPP server to the IANA registry is defined in the IANA Considerations section below. If the server uses both the IANA registry and DNS-based bootstrapping, then both methods MUST return the same location for the RPP server.

For DNS-based bootstrapping, an RPP server MUST publish an HTTPS resource record in a `_rpp` child zone for each DNS zone that is managed by the RPP server. The owner name of the HTTPS resource record MUST be the managed zone name itself, prepended with the `_rpp.` label (e.g., `_rpp.<zone>`).

If multiple HTTPS resource records are returned, the client MUST process them according to the priority rules defined in [RFC9460]. If multiple records have the same `SvcPriority`, the client SHOULD select one based on local policy.

The client MUST construct the URL for the well-known endpoint of the RPP server as follows, using the `TargetName` and port `SvcParams` from the HTTPS resource record:

- `https://<TargetName>:<port>/.well-known/rpp.json` when the port `SvcParam` is present and not 443
- `https://<TargetName>/.well-known/rpp.json` when the port `SvcParam` is absent or is 443

Example HTTPS resource record for an RPP server for the TLD "example" running HTTPS on port 443 at `rpp.example.`:

```
@ORIGIN example.  
_rpp IN HTTPS 1 . alpn=h2,h3
```

In this example, the well-known endpoint URL is `https://rpp.example/.well-known/rpp.json`.

Example HTTPS resource record for an RPP server using a different `TargetName` (`rpp-svr2.registry.example`) and a non-standard port (8443) for the TLD "example":

```
@ORIGIN example.  
_rpp IN HTTPS 1 rpp-svr2.registry.example. alpn=h2,h3 port=8443
```

In this example, the well-known endpoint URL is `https://rpp-svr2.registry.example:8443/.well-known/rpp.json`.

10. Discoverability

RPP server capabilities MUST be discoverable by clients. The server MUST provide a well-known endpoint at `/.well-known/rpp.json` at the root of the RPP server, this endpoint MUST return a JSON document containing the capabilities of the RPP server. The well-known endpoint MUST be accessible without authentication, and the client MUST be able to access this endpoint before authenticating with the server. The well-known endpoint MUST be accessible using the HTTP GET method and MUST return an HTTP status code 200 (OK) if the request was successful. The response message body MUST contain a JSON document describing the capabilities of the RPP server using the following fields:

- `base_url`: (required, string) The base URL for the RPP API, this is the URL that MUST be used as the base for all endpoint URL templates.
- `version`: (required, string) The version of the RPP API supported by the server, for example "1.0".
- `tlDs`: (required, array of strings) A list of TLDs supported by the server, for example "example", "org".

- **extensions:** (optional, array of extension objects) A list of supported extensions, each extension object **MUST** contain the following fields:
 - **name:** (required, string) A short name for the extension, for example "registry fee extension".
 - **id:** (required, string) A unique URN identifier for the extension, for example "urn:ietf:params:rpp:extension:registry-fee".
 - **version:** (required, string) The version of the extension supported by the server, for example "1.0".
 - **url:** (required, string) for standard extensions, this **MUST** be the URL for the extension in the IANA registry, for private extensions, this **MUST** be the URL for the extension specification, the domain name used in the URL **MUST** resolve and the URL **MUST** be accessible to the client.
- **profiles:** (optional, array of profile objects) A list of supported profiles, each profile object **MUST** contain the following fields:
 - **name:** (required, string) A short name for the profile, for example "domain provisioning profile".
 - **id:** (required, string) A unique URN identifier for the profile, for example "urn:ietf:params:rpp:profile:domain-provisioning-1.0".
 - **version:** (required, string) The version of the profile supported by the server, for example "1.0".
 - **url:** (required, string) for standard profiles, this **MUST** be the URL for the profile in the IANA registry, for private profiles, this **MUST** be the URL for the profile specification, the domain name used in the URL **MUST** resolve and the URL **MUST** be accessible to the client.
- **objects:** (required, array) A list of supported resource collections, for example "domains", "hosts", "entities".
- **authentication:** (optional, array of strings) A list of supported authentication methods, for example "Bearer", "Basic".
- **endpoints:** (required, array of endpoint objects) A list of available endpoints, each endpoint object **MUST** contain the following fields:
 - **name:** (required, string) A short name for the endpoint, for example "availability", "info", "poll", "create", "delete", "renewal" or "transfer".
 - **url_template:** (required, string) The URI template for the endpoint, using the syntax defined in [\[RFC6570\]](#).
- **maintenance:** (optional, array) An array containing information about upcoming planned maintenance windows of the server, with the following fields:
 - **start_time:** (required, string) The start time of the maintenance window in ISO 8601 format.
 - **end_time:** (required, string) The end time of the maintenance window in ISO 8601 format.
 - **description:** (optional, string) A human-readable description of the maintenance window.

The following template variables are defined for use in RPP endpoint URL templates. They are data object independent; the same variables are used regardless of which Data Object or Process Object the endpoint acts on.

- **collection:** The resource collection path segment, derived per Rule 1.
- **id:** The Unique Identifier value (as defined in [\[I-D.kowalik-rpp-data-objects\]](#)) of the resource instance within **collection**.
- **process-collection:** The process collection path segment, derived per Rule 3.
- **process-id:** The Unique Identifier value (as defined in [\[I-D.kowalik-rpp-data-objects\]](#)) of a specific process instance, scoped to its owner Data Object instance.

Example discovery response document:

```
{
  "base_url": "https://rpp.example/rpp/v1",
  "version": "1.0",
  "tlds": ["example", "org"],
  "extensions": [
    {
      "name": "RPP example extension",
      "id": "urn:ietf:params:rpp:extension:example-extension",
      "version": "1.0",
      "url": "https://www.iana.org/assignments/rpp-extensions/rpp-example-
extension-1.0"
    }
  ],
  "profiles": [
    {
      "name": "EPP compatibility profile",
      "id": "urn:ietf:params:rpp:profile:epp-compatibility",
      "version": "1.0",
      "url": "https://www.iana.org/assignments/rpp-profiles/epp-
compatibility-provisioning-profile-1.0"
    }
  ],
  "objects": ["domains", "hosts", "entities"],
  "endpoints": [
    {
      "name": "availability",
      "url_template": "/{collection}/{id}/availability"
    },
    {
      "name": "info",
      "url_template": "/{collection}/{id}"
    },
    {
      "name": "poll",
      "url_template": "/messages"
    },
    {
      "name": "create",
      "url_template": "/{collection}"
    }
  ],
  "authentication": ["Bearer"],
  "maintenance": [
    {
      "start_time": "2026-06-01T00:00:00Z",
      "end_time": "2026-06-01T06:00:00Z",
      "description": "Planned maintenance for server upgrades"
    }
  ]
}
```

10.1. Workflow

The steps for a typical workflow of provisioning an object using RPP without knowing the location and capabilities of the server are as follows, the first three steps are optional, the client can choose to skip any of these steps if it already has the required information from a previous interaction or configuration.

1. Bootstrap (optional): The client discovers the location of the RPP server by looking up the IANA registry for RPP servers or by performing a DNS SRV lookup as defined in [RFC2782].
2. Discover capabilities (optional): The client retrieves the capabilities of the RPP server by sending a GET request to the well-known endpoint at `/.well-known/rpp.json`.
3. Extract RPP URLs (optional): The client extracts the base URL and endpoint URL templates from the discovery response, and uses this information to construct the URLs for the desired operations.
4. Perform provisioning operations: The client performs provisioning operations by sending HTTP requests to the appropriate endpoint URLs, using the HTTP method and request message body as required by the specific operation.

11. Versioning

RPP is designed to be extensible and backward compatible. The version of the RPP API is indicated in the URL path, for example: `https://rpp.example/rpp/v1/`. The server **MUST** support at least one version of the RPP API, and **MUST** return a 404 Not Found status code for requests using an unsupported version. The versioning scheme uses the Semantic Versioning format defined in [SemVer], but only the major version number is used to indicate breaking changes. The minor and patch version numbers are not used in an URL path, but can be used in the media type or in the message body to indicate non-breaking changes.

The following RPP elements include versioning support:

- Endpoints: The server **MUST** support at least one version of the RPP API, and **MUST** return a 404 Not Found status code for requests using an unsupported version.
- Messages: A request and response message **MUST** include the version of the RPP API it is compatible with.
- Extensions: RPP extensions **MUST** include the version of the RPP API they are compatible with.
- Profiles: RPP profiles **MUST** include the version of the RPP API they are compatible with.
- Media types: RPP media types **MUST** include the version of the RPP API they are compatible with.
- Result codes: RPP result codes may be added by extensions or updates to the core RPP specification.

11.1. Endpoints

The `base_url` element of the RPP Discovery response **MAY** include the version of the RPP API supported by the server. The client **MUST** use this `base_url` for all subsequent requests to the server. For example, if the version is 1.2.3, the `base_url` is `https://rpp.example/rpp/v1/`, then the client **MUST** use this URL for all subsequent requests to the server, and **MUST** not use a different version in the URL path.

11.2. Messages

The `version` element of the RPP request and response messages **MUST** include the version of the RPP API that the message is compatible with. The server **MUST** reject requests with a version that is not supported by the server, and **MUST** return a RPP Client error code.

11.3. Extensions

The RPP server **MUST** include the version for each extension in the RPP Discovery response. The client **MUST** use this version information to determine which extensions are supported by the server, and to ensure that it uses the correct version of the extension when making requests to the server. A request using an extension **MUST** include the version of the extension. The server **MUST** reject requests using extensions with a version that is not supported by the server, and **MUST** return a RPP Client error code. The following is an example of how the version information for an extension can be included in the RPP Discovery response:

```
"extensions": [  
  {  
    "name": "RPP example extension",  
    "id": "urn:ietf:params:rpp:extension:example-extension",  
    "version": "1.0",  
    "url": "https://www.iana.org/assignments/rpp-extensions/rpp-example-  
extension-1.0"  
  }  
],
```

11.4. Profiles

The RPP server **MUST** include the version for each profile in the RPP Discovery response. The client **MUST** use this version information to determine which profiles are supported by the server, and to ensure that it uses the correct version of the profile when making requests to the server. A request using a profile **MUST** include the version of the profile. The server **MUST** reject requests using profiles with a version that is not supported by the server, and **MUST** return a RPP Client error code. The following is an example of how the version information for a profile can be included in the RPP Discovery response:

```
"profiles": [  
  {  
    "name": "EPP compatibility profile",  
    "id": "urn:ietf:params:rpp:profile:epp-compatibility",  
    "version": "1.0",  
    "url": "https://www.iana.org/assignments/rpp-profiles/epp-  
compatibility-provisioning-profile-1.0"  
  }  
]
```

12. Media types

RPP media types are used to indicate the format of the request and response messages, and MUST include a parameter indicating the name of the profile they are compatible with. The server MUST use the profile information in the media type to determine which features, extensions and versions to use when processing the request, and to ensure that it returns a response that is compatible with the client. The client MUST use the profile information in the media type to determine which features, extensions and versions to use when processing the response, and to ensure that it can correctly interpret the response.

The definition of profile parameters in media types is described in section

13. Profiles

A profile is a named set of protocol features and versions that are used to define the compatibility and capabilities of RPP server implementations, allowing for better interoperability between different implementations. Using profiles helps to simplify the implementation and deployment of RPP by providing a clear and concise way for the client and server to communicate their capabilities and requirements.

A profile is identified by a unique name and may be published as a standard profile in the IANA registry for RPP profiles, to promote interoperability and standardization across implementations. Standard profiles names MUST use the RPP URN namespace defined in this document, for example: `urn:ietf:params:rpp:profile:{profile_name}`. The profile name MUST be unique within the RPP URN namespace and SHOULD be descriptive of the features and versions included in the profile.

A profile can also be defined as a private profile, which is not published in the IANA registry, but is used by specific server implementations only. A private profile can be defined by a server operator to specify the features and versions supported by their implementation. If the profile is not published in the IANA registry, then the server operator MUST ensure that the profile name is globally unique to avoid conflicts with other profiles, the use of reverse domain name notation is RECOMMENDED for private profiles to ensure uniqueness.

13.1. Definition

A profile definition MUST contain the following fields, private profiles may contain additional fields as needed:

- `name`: A unique name that identifies the profile.
- `description`: A human-readable description of the profile and its intended use.
- `version`: The version of the profile.
- `rpp_version`: The minimum version of the RPP protocol that is supported or required for the profile.
- `objects`: A list of objects allowed for provisioning operations, for example "domains", "hosts", "entities", etc.
- `extensions`: A list of extensions, including their versions, that are supported or required for the profile.
- `profile`: A base profile that is extended by this profile. The base profile MUST be published in the IANA registry for RPP profiles, or be made available to the client using other means.

Example JSON representation for a standard profile named "example-profile" that supports RPP version 1.0 and includes two extensions, "rppExample" version 1.0 and "rppOther" version 1.1. The profile also "extends" the "base-profile" profile, which is defined in the IANA registry for RPP profiles and supports RPP version 1.0.

```
{
  "name": "urn:ietf:params:rpp:profile:example-profile",
  "description": "An example profile for provisioning objects using the RPP
protocol.",
  "version": "1.0",
  "rpp_version": "1.0",
  "objects": ["domains", "hosts", "entities"],
  "extensions": [
    {
      "rppExample": {
        "version": "1.0"
      },
      "rppOther": {
        "version": "1.1"
      }
    }
  ],
  "profile": {
    "name": "urn:ietf:params:rpp:profile:base-profile",
    "version": "1.0"
  }
}
```

13.2. Inheritance

A profile can include another profile, which is referred to as "inheriting" from that profile. When a profile inherits from another profile, it means that the features and versions defined in the inherited profile are also included in the inheriting profile. This allows for the creation of more complex profiles by building upon existing profiles. For example, a profile named "example-profile" could inherit from a base profile named "base-profile", which includes a set of common features and versions. The "example-profile" could then add additional features and versions on top of the ones defined in the "base-profile", or it can override some of the features and versions from the "base-profile". This allows for greater flexibility and modularity in defining profiles, as well as promoting reuse of common features and versions across different profiles. However, it is important to note that the use of inheritance in profiles can also make things more complicated, as it can create dependencies between profiles and make it harder to understand the features and versions included in a profile. Therefore, it is recommended to use inheritance with caution and to clearly document the relationships between profiles. The depth of inheritance MUST be limited to 1 to prevent excessive complexity, and profile designers MUST NOT create circular dependencies between profiles, where a profile inherits from itself directly.

This is an example of the parent base-profile used in the previous example, it contains a single extension "rppOther" version 1.0 and does not include any other profiles:

```
{
  "name": "urn:ietf:params:rpp:profile:base-profile",
  "description": "A base profile for provisioning objects using the RPP
protocol.",
  "version": "1.0",
  "rpp_version": "1.0",
  "objects": ["domains", "hosts", "entities"],
  "extensions": [
    {
      "rppOther": {
        "version": "1.0"
      }
    }
  ],
  "profile":
}
```

The example profile definition shown in [Section 13.1, Paragraph 3](#) uses the "base-profile" as its base and inherits its features and also overrides the version of the "rppOther" extension defined in the base profile.

13.3. Signalling

This document describes two distinct methods for signalling the profile used in a RPP request or response, the first method uses a dedicated HTTP header, the second method uses media type parameters. The two methods **MUST** not be used simultaneously in a single request or response. If both methods are used in a single request or response, then the server **MUST** return an HTTP error response and include a Problem Detail response in the message body.

13.3.1. Header signalling

The client **MUST** use the RPP-Profile header to indicate the name of the profile that is to be used for the request. The value of this header **MUST** be of the type parameter described in [\[RFC8941\]](#), the first parameter **MUST** uniquely identify a profile, for example `urn:ietf:params:rpp:profile:example-profile`, followed by the version parameter. If the server does not support the indicated profile or version, then the server **MUST** return an HTTP error response and include a Problem Detail response in the message body.

The ABNF for Profile header value is as follows:

```
profile-header = "profile" "=" profile-name ";" OWS "version" "=" version
profile-name   = token
version        = 1*DIGIT "." 1*DIGIT
```

Example:

```
RPP-Profile: profile=urn:ietf:params:rpp:profile:example-profile;version=1.0
```

13.3.2. Media type parameter signalling

When using Media type parameter signalling, the client and the server MUST use media type parameters in the Accept and Content-Type headers to indicate the name and version of the profile used in the request. The media type parameters MUST be defined as follows:

- **profile**: The value of this parameter MUST uniquely identify the profile, for example `urn:ietf:params:rpp:profile:example-profile`.
- **version**: The value of this parameter MUST indicate the version of the profile used in the request.

The ABNF for media type parameter signalling is as follows:

```
profile-parameter = "profile" "=" profile-name ";" OWS "version" "=" version
profile-name      = token
version           = 1*DIGIT "." 1*DIGIT
```

Example for the media type `application/rpp+json` with profile parameters indicating the use of the "example-profile" profile version 1.0.:

```
Accept: application/rpp+json; profile="urn:ietf:params:rpp:profile:example-
profile"; version="1.0"
Content-Type: application/rpp+json;
profile="urn:ietf:params:rpp:profile:example-profile"; version="1.0"
```

13.3.3. Process signalling

When a server creates a process resource as a side effect of a uniform interface operation, it signals this to the client using the Link response header [RFC8288] with `rel="rpp-process"`. The following target attributes are defined for this relation type:

- **process**: (REQUIRED) The process object identifier of the created process (e.g. `process="transferProcess"`).
- **processId**: (OPTIONAL) The server-assigned persistent identifier of the created process instance (e.g. `processId="XYZ-12345"`). If the server assigns a persistent identifier, this attribute MUST be included.
- **latest**: (OPTIONAL) The boolean value "true", indicating that the target URL is the latest process created of a given type. The target URL MAY use the "latest" mnemonic rather than a specific process identifier. There MUST NOT be more than one process of a given type with this attribute set to "true".

If multiple process objects are created, the server MUST include one Link header field per created process resource.

Example:

```
Link: <https://rpp.example/rpp/v1/domainNames/foo.example/processes/
transferProcesses/XYZ-12345>; rel="rpp-process"; process="transferProcess";
processId="XYZ-12345"; latest=true
```

14. Endpoints

Endpoints are described using URI Templates [RFC6570] relative to a discoverable base URL, as recommended by [RFC9205]. Some RPP endpoints do not require a request and/or response message.

14.1. HTTP Mapping Rules

All RPP endpoints are derived mechanically from the Data Object definitions in [I-D.ietf-rpp-data-objects]. No endpoint URL or HTTP method related to processing of provisioning objects is defined independently of a corresponding Data Object. The rules in this section MUST be applied to determine the URL path and HTTP method for any operation.

14.1.1. Rule 1: Collection Path Segment

Each Data Object has a stable, "Identifier" (e.g. "domainName", "contact", "host"). The URL path segment for a collection of such objects MUST be derived by applying the plural() function to the object's "Identifier". The "plural()" function appends an "s" or "es" to the identifier as per English language rules. Any irregular plural version MUST be defined in the data object specification.

```
{collection} = plural(dataObject.identifier)
```

Examples derived from current data object identifiers:

Data Object "Identifier"	"plural()" result	URL collection segment
"domainName"	"domainNames"	"/domainNames"
"contact"	"contacts"	"/contacts"
"host"	"hosts"	"/hosts"
"organisation"	"organisations"	"/organisations"

Table 2

14.1.2. Rule 2: Uniform Interface Operations

The four uniform interface operations defined in the RPP data object specification map to HTTP methods and URL paths as follows. "{collection}" is derived per Rule 1. "{id}" is the unique identifier value of the specific object instance.

A RPP client MAY use the HTTP GET method for informational requests only when no request data has to be added to the HTTP message body. Sending content using an HTTP GET request is discouraged in [RFC9110], there exist no generally defined semantics for

content received in a GET request. When an RPP operation requires additional input data, the client **MUST** use the HTTP POST, PUT or PATCH method and include any required data in the HTTP message body and HTTP headers.

TODO: the paragraph above looks like misplaced. Do we need it at all? The protocol defines if anything **MAY** be posted to the message body, so maybe this is a design consideration which does not belong to the final document?

Operation "Identifier"	HTTP Method	URL path
"create"	"POST"	"/{collection}"
"read"	"GET"	"/{collection}/{id}"
"update"	"PUT or PATCH"	"/{collection}/{id}"
"delete"	"DELETE"	"/{collection}/{id}"

Table 3

14.1.3. Rule 3: Direct Access Sub-Resource Path Segment

A data element whose "Direct Access" flag is set to true in its Data Object definition is additionally exposed as a sub-resource, nested under the URL of the resource instance that contains it. This rule applies uniformly and recursively: the containing resource instance **MAY** itself be a Direct Access sub-resource of a further-enclosing resource.

The fixed path segment for such a sub-resource **MUST** be derived by applying the `plural()` function to the Direct Access data element's own "Identifier" — not to the "Identifier" of the associated object type the element references.

```
{direct-access-segment} = plural(directAccessElement.identifier)
```

If the Direct Access element's cardinality is greater than 1, an individual associated object instance **MUST** be additionally addressed by appending its Unique Identifier value as a further path segment, "{unique-id}". [\[I-D.kowalik-rpp-data-objects\]](#) requires that any object type referenced by a Direct Access element of cardinality greater than 1 define a Unique Identifier for exactly this purpose. If the Direct Access element's cardinality is 0-1 or 1, no such extra segment is appended; the element's own path is already unambiguous.

```
{direct-access-path} = {container-path} "/" {direct-access-segment} [ "/"
{unique-id} ]
```

Applying Rule 3 recursively from the top-level Data Object down to every Direct Access element defined in [\[I-D.kowalik-rpp-data-objects\]](#) yields the following paths ("id" denotes the Unique Identifier value of the resource instance immediately to its left; it is instantiated per resource type as shown in the Derived Endpoint Reference below):

Container path	Direct Access element "Identifier"	Cardinality	Resulting path
"/domainNames/{id}"	"processes"	0-1	"/domainNames/{id}/processes"
"/contacts/{id}"	"processes"	0-1	"/contacts/{id}/processes"
"/hosts/{id}"	"processes"	0-1	"/hosts/{id}/processes"
"/organisations/{id}"	"processes"	0-1	"/organisations/{id}/processes"
"/organisations/{id}/users/{id}"	"processes"	0-1	"/organisations/{id}/users/{id}/processes"
"/{collection}/{id}/processes"	"transferProcess"	0+	"/{collection}/{id}/processes/transferProcesses/{process-id}"
"/{collection}/{id}/processes"	"renewProcess"	0+	"/{collection}/{id}/processes/renewProcesses/{process-id}"
"/{collection}/{id}/processes"	"restoreProcess"	0+	"/{collection}/{id}/processes/restoreProcesses/{process-id}"
"/{collection}/{id}/processes"	"createProcess"	0+	"/{collection}/{id}/processes/createProcesses/{process-id}"

Table 4

"{collection}" in the last four rows is derived per Rule 1 and stands for any owner Data Object's collection segment (e.g. "domainNames", "contacts"); the Processes element and its nested Process Object elements are defined once, generically, and Direct Access sub-resources derive identically regardless of the owner Data Object type. "{process-id}" is the Process Object's Unique Identifier value; Rule 4 defines the "latest" mnemonic as an additional way to address it.

The remaining rules in this section (Rules 4 through 6) apply Rule 3 specifically to Process Objects and describe how they are further exposed and interacted with.

14.1.4. Rule 4: Process Uniform Interface Operations

The four uniform interface operations from Rule 2 apply to Process Objects at the path derived per Rule 3. A server MAY assign a unique identifier "{process-id}" to each process instance and make it addressable by this identifier. The fixed keyword "latest" is used to address the most recent process instance when no specific process "{process-id}" is known or assigned.

Operation "Identifier"	HTTP Method	URL path
"create"	"POST"	"/{collection}/{id}/processes/{process-collection}"
"read"	"GET"	"/{collection}/{id}/processes/{process-collection}/latest"
"read" (specific instance)	"GET"	"/{collection}/{id}/processes/{process-collection}/{process-id}"
"delete"	"DELETE"	"/{collection}/{id}/processes/{process-collection}/latest"
"delete" (specific instance)	"DELETE"	"/{collection}/{id}/processes/{process-collection}/{process-id}"

Table 5

A started process MAY create a resource accessible using both the "latest" mnemonic and a server-assigned "{process-id}". If the server exposes any access to a process instances, access to the most recent instance via "latest" MUST be supported.

When a process is created and immediately completed by the server, a 201 Created response MAY still be provided with a "Location" header pointing to the created process resource.

If the server chooses not to expose any persistent process resource, it MUST return 200 OK instead of 201 Created.

The operations using the DELETE method MUST NOT accept a request body. Because the semantic of the HTTP DELETE method are not defined and may cause deployment issues. If a request body is present for any operation using the DELETE method, the server MUST reject the request with an appropriate error response.

This means that the "delete" operation (HTTP DELETE method) is not extensible, if the server needs to support additional data for a special "delete" operation, then a new process MUST be defined for this purpose.

14.1.5. Rule 5: Extended Process Operations

Operations on a Process Object beyond the uniform interface (e.g. "transferApprove", "transferReject", "report") are mapped to sub-resources of a specific process instance. The operation's "Identifier" is used unchanged as the final path segment. The HTTP method for all such extended operations MUST be POST.

Operation "Identifier"	HTTP Method	URL path
"{operationIdentifier}"	"POST"	"/{collection}/{id}/processes/{process-collection}/latest/{operationIdentifier}"
"{operationIdentifier}" (specific instance)	"POST"	"/{collection}/{id}/processes/{process-collection}/{process-id}/{operationIdentifier}"

Table 6

The operation "Identifier" MUST be used as-is with no transformation. The following examples are derived directly from the operation identifiers defined in the RPP data object specification:

Operation "Identifier"	Derived URL path (relative to "domains/foo.example")
"transferApprove"	"/domains/foo.example/processes/transferProcesses/latest/transferApprove"
"transferReject"	"/domains/foo.example/processes/transferProcesses/latest/transferReject"
"report"	"/domains/foo.example/processes/restoreProcesses/latest/report"

Table 7

14.1.6. Rule 6: Process Listing

A server MAY implement a listing facility for processes. If implemented, the following URL structures MUST be used by the client.

To retrieve all process instances of a given process type:

```
"GET /{collection}/{id}/processes/{process-collection}/"
```

To retrieve all process instances across all process types for an object:

```
"GET /{collection}/{id}/processes/"
```

A server MAY choose not to implement these endpoints, in which case it MUST return 404 Not Found or 501 Not Implemented.

14.2. Derived Endpoint Reference

The following table lists all current RPP endpoints, each derived by applying the rules above to the relevant data object and operation identifiers. The following table is non normative.

Operation	HTTP Method	URL path
Domain: read	"GET"	"/domainNames/{id}"
Domain: create	"POST"	"/domainNames"
Domain: update	"PUT or PATCH"	"/domainNames/{id}"
Domain: delete	"DELETE"	"/domainNames/{id}"
Contact: read	"GET"	"/contacts/{id}"
Contact: create	"POST"	"/contacts"
Contact: update	"PUT or PATCH"	"/contacts/{id}"

Operation	HTTP Method	URL path
Contact: delete	"DELETE"	"/contacts/{id}"
Host: read	"GET"	"/hosts/{id}"
Host: create	"POST"	"/hosts"
Host: update	"PUT or PATCH"	"/hosts/{id}"
Host: delete	"DELETE"	"/hosts/{id}"
Organisation: read	"GET"	"/organisations/{id}"
Organisation: create	"POST"	"/organisations"
Organisation: update	"PATCH"	"/organisations/{id}"
Organisation: delete	"DELETE"	"/organisations/{id}"
User: read	"GET"	"/organisations/{id}/users/{userId}"
User: create	"POST"	"/organisations/{id}/users"
User: update	"PATCH"	"/organisations/{id}/users/{userId}"
User: delete	"DELETE"	"/organisations/{id}/users/{userId}"
Transfer: create	"POST"	"/{collection}/{id}/processes/ transferProcesses"
Transfer: read	"GET"	"/{collection}/{id}/processes/ transferProcesses/latest"
Transfer: delete (cancel)	"DELETE"	"/{collection}/{id}/processes/ transferProcesses/latest"
Transfer: transferApprove	"POST"	"/{collection}/{id}/processes/ transferProcesses/latest/transferApprove"
Transfer: transferReject	"POST"	"/{collection}/{id}/processes/ transferProcesses/latest/transferReject"
Restore: create	"POST"	"/{collection}/{id}/processes/ restoreProcesses"
Restore: read	"GET"	"/{collection}/{id}/processes/ restoreProcesses/latest"
Restore: report	"POST"	"/{collection}/{id}/processes/ restoreProcesses/latest/report"
Renew: create	"POST"	"/{collection}/{id}/processes/ renewProcesses"

Operation	HTTP Method	URL path
Renew: read	"GET"	"/{collection}/{id}/processes/renewProcesses/latest"
Transfer: list	"GET"	"/{collection}/{id}/processes/transferProcesses"
Processes: list	"GET"	"/{collection}/{id}/processes"

Table 8

TODO: add availability and message queue

14.3. Availability for Creation

TODO: align with data object when availability is described there

The Availability for Creation endpoint is used to check whether an object can be successfully provisioned. Two distinct methods are defined for checking the availability of provisioning of an object, the first method uses the HEAD method for a quick check to find out if the object can be provisioned. The second method uses the GET method to retrieve additional information about the object's availability for provisioning, for example about pricing or additional requirements to be able to provision the requested object.

When the client uses the HTTP HEAD method, the server MUST respond with an HTTP status code 200 (OK) if the object can be provisioned or with an HTTP status code 404 (Not Found) if the object cannot be provisioned.

When the client uses the HTTP GET method, the server MUST respond with an HTTP status code 200 (OK) if the object can be provisioned. The server MUST include a message body containing more detailed availability information, for example about pricing or additional requirements to be able to provision the requested object. The message body MAY be an empty JSON object if no additional information is applicable.

If the object cannot be provisioned then the server MUST return an HTTP status code 404 (Not Found) and include a problem statement in the message body.

As an extension point the server MAY define and the client MAY use additional HTTP query parameters to further specify the check operation or the kind of response information that shall be returned. For example Registry Fee Extension [\[RFC8748\]](#) defines a possibility to request certain currency, only certain commands or periods. Such functionality would add query parameters, which could be used with GET request to receive additional pricing information with the response. HEAD request would not be affected in this case.

The server MUST respond with the same HTTP status code if the same URL is requested with HEAD and with GET.

- Request: HEAD|GET {collection}/{id}/availability
- Request message: None
- Response message: Optional availability response

Example request for a domain name that is not available for provisioning:

```
HEAD domains/foo.example/availability HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example response:

```
HTTP/2 404 Not Found
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
RPP-Cltrid: ABC-12345
RPP-Svtrid: XYZ-12345
RPP-code: 01000
Content-Length: 0
```

14.4. Resource Information

The Object Info request MUST use the HTTP GET method on a resource identifying an object instance (Rule 2, read operation). If the object has authorization information attached then the client MUST use an empty message body and include the RPP-Authorization HTTP header. If the authorization is linked to a database object the client MUST also include the roid in the RPP-Authorization header. The client MAY also use a message body that includes the authorization information, the client MUST then not use the RPP-Authorization header.

Example request for an object not using authorization information:

```
GET /rpp/v1/domainNames/foo.example HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example request using RPP-Authorization header for an object that has attached authorization information:

```
GET /rpp/v1/domainNames/foo.example HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
RPP-Authorization: authinfo value=TXkgU2VjcmV0IFRva2Vu
```

Example Info response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 424
Content-Type: application/rpp+json
Content-Language: en
RPP-code: 01000

TODO: JSON message here
```

14.5. Create Resource

The client **MUST** use the HTTP POST method on a resource identifying a collection of object instances (Rule 2, create operation).

Example Domain Create request for a new domain name `foo.example`:

```
POST /rpp/v1/domainNames HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Content-Type: application/rpp+json
Accept-Language: en
Content-Length: 220

TODO
```

Example Domain Create response for a new domain name `foo.example`:

```
HTTP/2 201 Created
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Language: en
Content-Length: 642
Content-Type: application/rpp+json
Location: https://rpp.example/rpp/v1/domainNames/foo.example
RPP-code: 01000
```

TODO

Example Domain Create response where a createProcess object was implicitly created:

```
HTTP/2 201 Created
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Language: en
Content-Type: application/rpp+json
Location: https://rpp.example/rpp/v1/domainNames/foo.example
Link: <https://rpp.example/rpp/v1/domainNames/foo.example/processes/
createProcesses/latest>; rel="rpp-process"
RPP-code: 01000
```

TODO

14.6. Delete Resource

The client MUST use the HTTP DELETE method on a resource identifying a unique object instance (Rule 2, delete operation).

Example Domain Delete request:

```
DELETE /rpp/v1/domainNames/foo.example HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example Domain Delete response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 80
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
RPP-code: 01000
```

TODO

14.7. Update Resource

RPP supports two complementary update operations for modifying an existing object instance, each with its own semantics and use cases:

- **Full update** (HTTP PUT): The client sends a complete replacement representation of the object. The server **MUST** replace the stored object with the provided representation. Any attributes not present in the request body **MUST** be treated as absent and cleared or reset to their default values, subject to server policy. The client **MUST** send all read-write attributes required by the data model, not just the changed ones. The client **MUST** not send any create-only attributes.
- **Partial update** (HTTP PATCH): The client sends only the attributes to be modified. The server **MUST** apply only the changes indicated in the request body and leave all other attributes unchanged. Data representation of the partial update payload determines how the changes are transmitted between client and server and applied to the data object. The client **MAY** send changes to any read-write attributes defined in the data model. The client **MUST** not send any create-only attributes.

Both operations **MUST** be performed on a URL identifying a unique object instance (Rule 2). The request body **MUST** contain a valid object representation in the negotiated media type.

The server **MUST** respond with HTTP status code 200 (OK) and include the updated object representation in the response body.

Example full update request (PUT):

```
PUT /rpp/v1/domainNames/foo.example HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Content-Type: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
Content-Length: 252
```

TODO

Example full update response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 80
Content-Type: application/rpp+json
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
RPP-code: 01000
```

TODO

Example partial update request (PATCH):

```
PATCH /rpp/v1/domainNames/foo.example HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Content-Type: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
Content-Length: 252
```

TODO

Example partial update response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 80
Content-Type: application/rpp+json
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
RPP-code: 01000
```

TODO

14.8. Processes

Each provisioning object may be related to one or more running processes, such as a transfer or renewal. Each process has its own data, distinct from the data of the provisioning object itself, and may be interacted with using its own set of operations.

All process resources **MUST** exist under the `/collection/{id}/processes/{process-collection}` path, where `{process-collection}` is derived per Rule 3.

14.8.1. Relation to object representation

A uniform interface operation MAY require process data in addition to the object representation data. How the process data is embedded in the request body MUST be defined in the corresponding representation specification.

14.8.2. Response with information about created process

When the server creates a process object as a side effect of the operation, it MUST signal this to the client as described in [Section 13.3.3](#).

14.8.3. Restore Resource

TODO: this needs update once restoreProcess is defined in Data Objects

14.8.4. Renew Resource

Renew is modelled as a Process Object with its own lifecycle. The `renewProcess` object identifier yields the `renewProcesses` collection segment per Rule 3.

The client MUST use the HTTP POST method to create a new renew process (Rule 4, create).

Not every object resource includes support for the renew command. The response MUST include the Location header for the created renew process resource.

Example Domain Renew request:

```
POST /rpp/v1/domainNames/foo.example/processes/renewProcesses HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Content-Type: application/rpp+json
RPP-Cltrid: ABC-12345
Accept-Language: en
Content-Length: 96

{
  "expiryDate": "2025-09-08",
  "renewalPeriod": {
    "unit": "y",
    "value": 1
  }
}
```

Example Renew response:


```
HTTP/2 201 Created
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Language: en
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
Content-Length: 85
Location: https://rpp.example/rpp/v1/domainNames/foo.example/processes/
renewProcesses/XYZ-12345
Content-Type: application/rpp+json
RPP-code: 01000

{
  "expiryDate": "2026-09-08"
}
```

14.8.5. Transfer Resource

The Transfer operation manages the change of sponsoring client for a provisioned object. Transfer is modelled as a Process Object with its own lifecycle. The `transferProcess` object identifier yields the `transferProcesses` collection segment per Rule 3.

14.8.5.1. Start

The initiating client MUST use the HTTP POST method to create a new transfer process (Rule 4, create).

Example request not using object authorization:

```
POST /rpp/v1/domainNames/foo.example/processes/transferProcesses HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
Content-Length: 320

{
  "transferDir": "push",
  "gainingClientId": "ClientX"
}
```

Example request using object authorization:

```
POST /rpp/v1/domainNames/foo.example/processes/transferProcesses HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
RPP-Cltrid: ABC-12345
RPP-Authorization: authinfo value=TXkgU2VjcmV0IFRva2Vu
Accept-Language: en
Content-Length: 320

{
  "transferDir": "pull"
}
```

Example Transfer response:

```
HTTP/2 201 Created
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Language: en
Content-Length: 182
Content-Type: application/rpp+json
Location: https://rpp.example/rpp/v1/domainNames/foo.example/processes/
transferProcesses/latest
RPP-code: 01001

{
  "trStatus": "pending",
  "reqClientId": "ClientX",
  "actClientId": "ClientY",
  "requestDate": "2000-06-06T22:00:00.0Z",
  "actionDate": "2000-06-11T22:00:00.0Z",
  "exDate": "2002-09-08T22:00:00.0Z"
}
```

14.8.5.2. Status

A transfer process resource may not exist when no transfer has been initiated for the specified object. The client **MUST** use the HTTP GET method and **MUST NOT** add content to the HTTP message body.

Example domain name Transfer Status request:

```
GET /rpp/v1/domainNames/foo.example/processes/transferProcesses/latest HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example Transfer Query response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 230
Content-Type: application/rpp+json
Content-Language: en
RPP-code: 01000

{
  "trStatus": "pending",
  "reqClientId": "ClientX",
  "actClientId": "ClientY",
  "requestDate": "2000-06-06T22:00:00.0Z",
  "actionDate": "2000-06-11T22:00:00.0Z",
  "exDate": "2002-09-08T22:00:00.0Z"
}
```

14.8.5.3. Cancel

The initiating client cancels its pending transfer request using the DELETE method on the process resource (Rule 4, delete).

Example request:

```
DELETE /rpp/v1/domainNames/foo.example/processes/transferProcesses/latest
HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 80
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
RPP-code: 01000

{
  "trStatus": "clientCancelled",
  "reqClientId": "ClientX",
  "actClientId": "ClientY",
  "requestDate": "2000-06-06T22:00:00.0Z",
  "actionDate": "2000-06-11T22:00:00.0Z"
}
```

14.8.5.4. Reject

The currently sponsoring client rejects a pending transfer. This is an extended process operation; the operation identifier `transferReject` is used unchanged as the final path segment (Rule 5).

Example request:

```
POST /rpp/v1/domainNames/foo.example/processes/transferProcesses/latest/
transferReject HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 80
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
RPP-code: 01000

{
  "trStatus": "clientRejected",
  "reqClientId": "ClientX",
  "actClientId": "ClientY",
  "requestDate": "2000-06-06T22:00:00.0Z",
  "actionDate": "2000-06-11T22:00:00.0Z"
}
```

14.8.5.5. Approve

The currently sponsoring client approves a pending transfer. The operation identifier `transferApprove` is used unchanged as the final path segment (Rule 5).

Example request:

```
POST /rpp/v1/domainNames/foo.example/processes/transferProcesses/latest/
transferApprove HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
Content-Length: 0
```

Example response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 80
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
RPP-code: 01000

{
  "trStatus": "clientApproved",
  "reqClientId": "ClientX",
  "actClientId": "ClientY",
  "requestDate": "2000-06-06T22:00:00.0Z",
  "actionDate": "2000-06-11T22:00:00.0Z"
}
```

14.9. Messages

14.9.1. Retrieve

TODO: update when covered in data objects

The messages endpoint is used for retrieving messages stored on the server for the client to process.

- Request: GET /messages
- Request message: None
- Response message: Poll response

The client MUST use the HTTP GET method on the messages resource collection to request the message at the head of the queue.

Example request:

```
GET messages HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Length: 312
Content-Type: application/rpp+json
Content-Language: en
RPP-code: 01301
```

TODO

14.9.2. Delete

TODO: update when covered in data objects

- Request: DELETE /messages/{id}
- Request message: None
- Response message: Poll Ack response

The client MUST use the HTTP DELETE method to acknowledge receipt of a message from the queue. The "msgID" attribute of a received RPP Poll message MUST be included in the message resource URL, using the {id} path element. The server MUST use RPP headers to return the RPP result code and the number of messages left in the queue. The server MUST NOT add content to the HTTP message body of a successful response, the server may add content to the message body of an error response.

Example request:

```
DELETE messages/12345 HTTP/2
Host: rpp.example
Authorization: Bearer <token>
Accept: application/rpp+json
Accept-Language: en
RPP-Cltrid: ABC-12345
```

Example response:

```
HTTP/2 200 OK
Date: Wed, 24 Jan 2024 12:00:00 UTC
Server: Example RPP server v1.0
Content-Language: en
RPP-code: 01000
RPP-Queue-Size: 0
RPP-Svtrid: XYZ-12345
RPP-Cltrid: ABC-12345
Content-Length: 145
```

TODO

15. Extension Framework

TODO

16. RPP Result Codes

RPP result codes are used to indicate the result of an RPP request. They are returned in the RPP-Code header of the HTTP response. The format of the RPP result code is a 5-digit string, where the first digit **MUST** always be "1", the second digit indicates the class of the result, and the remaining four digits indicate the specific result within that class, his allows implementers to define more specific result codes within each class. The classes of RPP result codes are designed to match the classes of HTTP status codes, to facilitate mapping between RPP result codes and HTTP status codes. The classes of RPP result codes are defined as follows:

- 11xxx: Informational
- 12xxx: Success
- 13xxx: Reserved for future use
- 14xxx: Client error
- 15xxx: Server error

The following RPP result codes are defined and used in this document:

RPP Result Code	HTTP Status Code	Description
12000	200 OK	Command completed successfully
12001	201 Created	Command completed successfully and a new resource was created

Table 9

17. Authentication and Authorization

Due to the stateless nature of RPP, the client **MUST** include the authentication credentials in each HTTP request. This **MAY** be done by using JSON Web Tokens (JWT) [RFC7519] or Basic authentication [RFC7617]. The server **MUST** validate the authentication credentials on each request and reject any request with invalid credentials with an appropriate HTTP status code.

When using JWTs for OAuth 2.0 [RFC6749] Access Tokens, the JWT profile described in [RFC9068] **MUST** be used. It is **RECOMMENDED** to use short-lived tokens and to implement token revocation mechanisms to mitigate the risk of token compromise. If sensitive information is included in the JWT payload, it **MUST** be encrypted to prevent unauthorized access when the token is persistent to a storage device. Furthermore, the best practices for JWT usage as outlined in [RFC8725] **MUST** be followed.

18. IANA Considerations

18.1. URN Sub-namespace for RPP (urn:ietf:params:rpp)

The IANA is requested to add the following value to the "IETF URN Sub-namespace for Registered Protocol Parameter Identifiers" registry, following the template in [\[RFC3553\]](#):
TODO: add filled in template, if we decide to use URN for profile identification, for example "urn:ietf:params:rpp:profile:example-profile"

```
Registry name: rpp
Specification: This Document
Repository: ?
Index value: ?
```

18.2. RPP registry group

The IANA is requested to create a new registry group for RPP, this will be used to group together all the RPP-related registries such as those for discovery URLs, extensions and profiles.

```
Name of the registry: RESTful Provisioning Protocol (RPP)
Reference: This Document
IANA Registry Reference: TODO
```

18.3. RPP Discovery registry

The IANA is requested to create a new registry for RPP discovery URLs, this registry will be used to register the well-known URLs for RPP discovery endpoints, used by RPP clients to discover the capabilities of a RPP server.

```
Name of the registry: RPP Discovery URLs
Registry group: RESTful Provisioning Protocol (RPP)
Registration procedure: Expert Review
```

Fields to be registered:

- `tld`: The top-level domain (TLD) for which the discovery URL is applicable, for example "example".
- `url`: The URL for the discovery endpoint, for example "<https://rpp.example/.well-known/rpp.json>".
- `description`: A human-readable description of the discovery URL and its intended use.

18.4. RPP Extension registry

The IANA is requested to create a new registry for RPP extensions, this registry will be used to register standardized extensions to the RPP protocol. Extensions are defined as additional features or capabilities that can be added to the base RPP protocol, for example support for additional resource types, additional operations or additional authentication methods.

```
Name of the registry: RPP Extensions
Registry group: RESTful Provisioning Protocol (RPP)
Registration procedure: Expert Review
```

Fields to be registered:

- **name:** The name of the extension, for example "RPP example extension".
- **version:** The version of the extension, for example "1.0".
- **url:** The URL for the extension specification, for example "<https://www.iana.org/assignments/rpp-extensions/rpp-example-extension-1.0>".
- **description:** A human-readable description of the extension and its intended use.

18.5. RPP Profile registry

The IANA is requested to create a new registry for RPP profiles, this registry will be used to register standardized profiles for the RPP protocol. Profiles are defined as specific configurations of the RPP protocol that are designed to meet the needs of specific use cases or environments, for example an EPP compatibility profile that defines a set of RPP features and behaviors that are compatible with the Extensible Provisioning Protocol (EPP) [RFC5730].

```
Name of the registry: RPP Profiles
Registry group: RESTful Provisioning Protocol (RPP)
Registration procedure: Expert Review
```

Fields to be registered:

- **name:** The name of the profile, for example "EPP compatibility profile".
- **id:** A unique URN identifier for the profile, for example "urn:ietf:params:rpp:profile:epp-compatibility-1.0".
- **version:** The version of the profile, for example "1.0".
- **url:** The URL for the profile specification, for example "<https://www.iana.org/assignments/rpp-profiles/epp-compatibility-provisioning-profile-1.0>".
- **description:** A human-readable description of the profile and its intended use.

18.6. RPP Result Codes Registry

The IANA is requested to create a new registry "RPP Result codes", this registry will be used to register RPP result codes defined in this document and in future RPP specifications and extensions.

Name of the registry: RPP Result codes
Registry group: RESTful Provisioning Protocol (RPP)
Registration procedure: Expert Review

Fields to be registered:

- code: The RPP result code, for example "12000".
- description: A human-readable description of the result code and its intended use.

18.7. Link Relation Type: rpp-process

The IANA is requested to register the following link relation type in the "Link Relation Types" registry, following the template in [\[RFC8288\]](#):

Relation Name: rpp-process
Description: Identifies a process resource that was implicitly created as a side effect of a uniform interface operation on the target resource. The context resource is the provisioning object on which the operation was performed; the target resource is the created process instance. The following target attributes are defined for this relation type: "process" (the Process Object Identifier, REQUIRED), "processId" (the server-assigned instance identifier, OPTIONAL), and "latest" (the boolean value "true", OPTIONAL, indicating the target URL uses the "latest" mnemonic).
Reference: This document

18.8. RPP Media Type (application/rpp+json)

The IANA is requested to add the following RPP media type to the "Media Types" registry, following the template in [\[RFC6838\]](#):

```
Type name: application
Subtype name: rpp+json
Required parameters: version
Optional parameters: "N/A"
Encoding considerations: "N/A"
Security considerations: "N/A"
Interoperability considerations: "N/A"
Published specification: This document
Applications that use this media type: RPP protocol and extensions
Fragment identifier considerations: "N/A"
Additional information: "N/A"
Person & email address to contact for further information: Author's email
address
Intended usage: COMMON
Restrictions on usage: "N/A"
Author: Document authors
Change controller: Document authors
Provisional registration: No
```

19. Internationalization Considerations

TODO

20. Security Considerations

RPP relies on the security of the underlying HTTP transport, hence the best common practices for securing HTTP described in [\[RFC9325\]](#) also apply to RPP and MUST be followed by RPP implementations.

Data confidentiality and integrity MUST be enforced. Every client and server interaction MUST be encrypted using TLS version 1.3 [\[RFC8446\]](#). Future versions of TLS MAY be used as they become available and are deemed secure.

21. Change History

21.1. Version draft-wullink-rpp-core-05 to draft-ietf-rpp-core-00

- Renamed the document name to "draft-ietf-rpp-core" and reset version number to 00 (Issue #80)
- Added Organisation and User resource type, with create, read, update and delete operations. (Issue #67)
- Added section about Mapping to EPP. (Issue #55)
- Described full and partial update operations, using HTTP PUT and PATCH methods respectively. (Issue #21)
- Added support for embedding process data in uniform interface operations. Registered the rpp-process link relation type with IANA.
- Generalised process sub-resource URL derivation into Rule 3 (Direct Access Sub-Resource Path Segment); Rules 4 through 6 now apply this general rule to Process Objects instead of deriving process paths independently.

- Corrected the Renew process collection segment from `renewalProcesses` to `renewProcesses` in the Renew Resource section and the Derived Endpoint Reference table, matching the `renewProcess` identifier in [\[I-D.kowalik-rpp-data-objects\]](#).
- Aligned the URL template variables in Discoverability with Unique Identifier terminology from [\[I-D.kowalik-rpp-data-objects\]](#) and with the `{process-collection}`/`{process-id}` naming used in the Endpoints section; removed the unused `process_name` variable.

21.2. Version 04 to 05

- Added Bootstrap and Discovery sections to the document, describing how a client can discover the location and capabilities of an RPP server
- Added IANA Considerations section with a request for new RPP discovery URLs, extensions and profile URLs.

21.3. Version 03 to 04

- Added a new section on versioning, describing how versioning is applied to different RPP elements. (Issue #39)
- Added IANA request for RPP Result codes registry, and added a table with example RPP result codes. (Issue #39)
- Added IANA request for URN sub-namespace for RPP. (Issue #39)
- Added a new "Profiles" section to describe how to define and use profiles. (Issue #43)
- Added a new section "Authentication and Authorization". (Issue #37)
- Updated the "Security Considerations" section to include transport security. (Issue #37)
- Added IANA registration request for the new RPP media type. (Issue #40)

21.4. Version 02 to 03

- Added use of Problem Detail [\[RFC9457\]](#) for error responses

21.5. Version 01 to 02

- Updated the examples, changed from `".example.org"` to `".example"`
- Merged the RPP-EPP-Code and RPP-Code headers into a single RPP-Code header
- Update the RPP-Authorization header to match the HTTP Authorization header format
- Added new process path segment and process representations
- Updated the Check request to now use an "availability" path segment and support both GET and HEAD methods

21.6. Version 00 to 01

- Updated "Request Headers" and "Response Headers" section
- Changed transfer resource URL and HTTP method for reject, approve and cancel, in order to make the API easier to use

21.7. Version 00 (draft-rpp-core) to 00 (draft-wullink-rpp-core)

- Renamed the document name to "draft-wullink-rpp-core"
- Removed sections: Design Considerations, Resource Naming Convention, Session Management, HTTP Layer, Content Negotiation, Object Filtering, Error Handling
- Renamed Commands section to Endpoints
- Removed text about extensions

- Changed naming to be less EPP like and more RDAP like

22. Acknowledgements

The authors would like to thank the following people for their helpful text contributions, comments and suggestions.

- Q Misell, AS207960 Cyfyngedig

23. References

23.1. Normative References

- [I-D.ietf-rpp-data-objects]** Kowalik, P. and M. Wullink, "RESTful Provisioning Protocol (RPP) Data Objects", Work in Progress, Internet-Draft, draft-ietf-rpp-data-objects-01, 3 July 2026, <<https://datatracker.ietf.org/doc/html/draft-ietf-rpp-data-objects-01>>.
- [I-D.kowalik-rpp-data-objects]** Kowalik, P. and M. Wullink, "RESTful Provisioning Protocol (RPP) Data Objects", Work in Progress, Internet-Draft, draft-kowalik-rpp-data-objects-04, 12 June 2026, <<https://datatracker.ietf.org/doc/html/draft-kowalik-rpp-data-objects-04>>.
- [REST]** Fielding, R., "Architectural Styles and the Design of Network-based Software Architectures", 2000, <http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm>.
- [RFC2119]** Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997, <<https://www.rfc-editor.org/info/rfc2119>>.
- [RFC2616]** Fielding, R., Gettys, J., Mogul, J., Frystyk, H., Masinter, L., Leach, P., and T. Berners-Lee, "Hypertext Transfer Protocol -- HTTP/1.1", RFC 2616, DOI 10.17487/RFC2616, June 1999, <<https://www.rfc-editor.org/info/rfc2616>>.
- [RFC2782]** Gulbrandsen, A., Vixie, P., and L. Esibov, "A DNS RR for specifying the location of services (DNS SRV)", RFC 2782, DOI 10.17487/RFC2782, February 2000, <<https://www.rfc-editor.org/info/rfc2782>>.
- [RFC3553]** Mealling, M., Masinter, L., Hardie, T., and G. Klyne, "An IETF URN Subnamespace for Registered Protocol Parameters", BCP 73, RFC 3553, DOI 10.17487/RFC3553, June 2003, <<https://www.rfc-editor.org/info/rfc3553>>.
- [RFC3986]** Berners-Lee, T., Fielding, R., and L. Masinter, "Uniform Resource Identifier (URI): Generic Syntax", STD 66, RFC 3986, DOI 10.17487/RFC3986, January 2005, <<https://www.rfc-editor.org/info/rfc3986>>.
- [RFC5730]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP)", STD 69, RFC 5730, DOI 10.17487/RFC5730, August 2009, <<https://www.rfc-editor.org/info/rfc5730>>.
- [RFC5731]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP) Domain Name Mapping", STD 69, RFC 5731, DOI 10.17487/RFC5731, August 2009, <<https://www.rfc-editor.org/info/rfc5731>>.
- [RFC5733]** Hollenbeck, S., "Extensible Provisioning Protocol (EPP) Contact Mapping", STD 69, RFC 5733, DOI 10.17487/RFC5733, August 2009, <<https://www.rfc-editor.org/info/rfc5733>>.

-
- [RFC6570] Gregorio, J., Fielding, R., Hadley, M., Nottingham, M., and D. Orchard, "URI Template", RFC 6570, DOI 10.17487/RFC6570, March 2012, <<https://www.rfc-editor.org/info/rfc6570>>.
- [RFC6648] Saint-Andre, P., Crocker, D., and M. Nottingham, "Deprecating the "X-" Prefix and Similar Constructs in Application Protocols", BCP 178, RFC 6648, DOI 10.17487/RFC6648, June 2012, <<https://www.rfc-editor.org/info/rfc6648>>.
- [RFC6749] Hardt, D., Ed., "The OAuth 2.0 Authorization Framework", RFC 6749, DOI 10.17487/RFC6749, October 2012, <<https://www.rfc-editor.org/info/rfc6749>>.
- [RFC6838] Freed, N., Klensin, J., and T. Hansen, "Media Type Specifications and Registration Procedures", BCP 13, RFC 6838, DOI 10.17487/RFC6838, January 2013, <<https://www.rfc-editor.org/info/rfc6838>>.
- [RFC7519] Jones, M., Bradley, J., and N. Sakimura, "JSON Web Token (JWT)", RFC 7519, DOI 10.17487/RFC7519, May 2015, <<https://www.rfc-editor.org/info/rfc7519>>.
- [RFC7617] Reschke, J., "The 'Basic' HTTP Authentication Scheme", RFC 7617, DOI 10.17487/RFC7617, September 2015, <<https://www.rfc-editor.org/info/rfc7617>>.
- [RFC8288] Nottingham, M., "Web Linking", RFC 8288, DOI 10.17487/RFC8288, October 2017, <<https://www.rfc-editor.org/info/rfc8288>>.
- [RFC8446] Rescorla, E., "The Transport Layer Security (TLS) Protocol Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018, <<https://www.rfc-editor.org/info/rfc8446>>.
- [RFC8725] Sheffer, Y., Hardt, D., and M. Jones, "JSON Web Token Best Current Practices", BCP 225, RFC 8725, DOI 10.17487/RFC8725, February 2020, <<https://www.rfc-editor.org/info/rfc8725>>.
- [RFC8941] Nottingham, M. and P. Kamp, "Structured Field Values for HTTP", RFC 8941, DOI 10.17487/RFC8941, February 2021, <<https://www.rfc-editor.org/info/rfc8941>>.
- [RFC9068] Bertocci, V., "JSON Web Token (JWT) Profile for OAuth 2.0 Access Tokens", RFC 9068, DOI 10.17487/RFC9068, October 2021, <<https://www.rfc-editor.org/info/rfc9068>>.
- [RFC9110] Fielding, R., Ed., Nottingham, M., Ed., and J. Reschke, Ed., "HTTP Semantics", STD 97, RFC 9110, DOI 10.17487/RFC9110, June 2022, <<https://www.rfc-editor.org/info/rfc9110>>.
- [RFC9205] Nottingham, M., "Building Protocols with HTTP", BCP 56, RFC 9205, DOI 10.17487/RFC9205, June 2022, <<https://www.rfc-editor.org/info/rfc9205>>.
- [RFC9325] Sheffer, Y., Saint-Andre, P., and T. Fossati, "Recommendations for Secure Use of Transport Layer Security (TLS) and Datagram Transport Layer Security (DTLS)", BCP 195, RFC 9325, DOI 10.17487/RFC9325, November 2022, <<https://www.rfc-editor.org/info/rfc9325>>.
- [RFC9457] Nottingham, M., Wilde, E., and S. Dalal, "Problem Details for HTTP APIs", RFC 9457, DOI 10.17487/RFC9457, July 2023, <<https://www.rfc-editor.org/info/rfc9457>>.
- [RFC9460] Schwartz, B., Bishop, M., and E. Nygren, "Service Binding and Parameter Specification via the DNS (SVCB and HTTPS Resource Records)", RFC 9460, DOI 10.17487/RFC9460, November 2023, <<https://www.rfc-editor.org/info/rfc9460>>.
-

[RFC9535] Gössner, S., Ed., Normington, G., Ed., and C. Bormann, Ed., "JSONPath: Query Expressions for JSON", RFC 9535, DOI 10.17487/RFC9535, February 2024, <<https://www.rfc-editor.org/info/rfc9535>>.

[SemVer] Semantic Versioning, "Semantic Versioning 2.0.0", <<https://semver.org/>>.

23.2. Informative References

[RFC8748] Carney, R., Brown, G., and J. Frakes, "Registry Fee Extension for the Extensible Provisioning Protocol (EPP)", RFC 8748, DOI 10.17487/RFC8748, March 2020, <<https://www.rfc-editor.org/info/rfc8748>>.

Authors' Addresses

Maarten Wullink

SIDN Labs

Email: maarten.wullink@sidn.nl

URI: <https://sidn.nl/>

Pawel Kowalik

DENIC

Email: pawel.kowalik@denic.de

URI: <https://denic.de/>